

# 432 Class 10 Slides

[thomaseLove.github.io/432](https://thomaseLove.github.io/432)

2022-02-10

# Today's Agenda

## The tidymodels framework

- Using tidymodels tools to develop a linear regression model
  - Pre-processing activities
  - Model building (with multiple fitting engines)
  - Measuring model effectiveness
  - Creating a model workflow

## Next Time (Class 11)

- Using tidymodels tools to develop a logistic regression model

# Setup

```
library(here); library(conflicted)
library(knitr); library(magrittr); library(janitor)

library(tidymodels)
library(tidyverse)

theme_set(theme_bw())

conflict_prefer("select", "dplyr")
conflict_prefer("filter", "dplyr")
```

# Regression Frameworks

Generally, regression allows us to summarize how predictions (or average values) of an outcome vary across individuals defined by a set of predictors. Some of the most important uses of regression are:

- **Prediction**, which involves both modeling existing observations and forecasting new data.
- **Exploring Associations**, where we summarize how well a set of variables predicts the outcome.
- **Extrapolation**, where we are adjusting for known differences between the observed sample of data and a population of interest.
- **Causal Inference**, where we are estimating the effect of a treatment, by comparing outcomes under treatment or control, or under different levels of a treatment<sup>1</sup>.

Source: Gelman, Hill and Vehtari, *Regression and Other Stories*

---

<sup>1</sup>My 500 course spends a whole semester on one important part of this subject.

# Research Questions for Regression Models

- “How effectively can [insert quantitative outcome] be predicted using [insert predictor(s)]?” for a linear regression project, and
- “How effectively can [insert binary outcome] be predicted using [insert predictor(s)]?” for a logistic regression project.

If you're struggling with this, or if your research question isn't in the form of a question, consider these approaches. Advantages:

- 1 regression can help provide an answer to these questions and in discussing your results you'll need to answer the questions
- 2 framing models in terms of exploring associations has some value for the tools we're discussing and
- 3 it's pretty clear what you're doing, based just on your research question.

If you're doing something else, I still need to think that you meet standards (1) and (3) at least.

# Using R to fit Regression Models

For linear models, we have:

- `lm` to fit models for quantitative outcomes, compute and plot predictions and residuals, obtain confidence intervals, etc.
- `ols` from the `rms` package to save and explore additional components of the model's fit and to (slightly) expand the capacity for `lm` fits to incorporate non-linear terms and multiple imputations.

For logistic models, we have:

- `glm` to fit models for binary outcomes, compute and plot predictions, hypothesis tests and confidence intervals
- `lrm` from `rms` to save and explore additional components of the model's fit and to (slightly) expand the capacity for `lm` fits to incorporate non-linear terms and multiple imputations.

These are by no means the only options for fitting or working with models.

# What are tidymodels?

The `tidymodels` collection of packages in R use tidyverse principles to facilitate modeling and machine learning work. The key idea is to develop a consistent framework for modeling, including:

- pre-processing data, which includes identifying variables and their roles, re-expression of outcomes, creation of features (predictors)
- building a model (potentially with multiple fitting “engines”)
- developing a re-usable workflow
- evaluating the fit of one model or various models with a variety of validation strategies

Visit the `tidymodels` website at <https://www.tidymodels.org/>.

# Core Tidymodels Packages

Install many of the packages in the tidymodels ecosystem with `install.packages(tidymodels)`.

When you use `library(tidymodels)`, this makes the core packages available in your R session. They include:

- `rsample` which will help with data splitting and resampling
- `parsnip` which provides a tidy, unified interface for models
- `recipes` for data pre-processing and feature engineering
- `yardstick` for measuring model effectiveness
- `broom` for converting R objects into predictable formats
- `workflows` for bundling together pre-processing, modeling and post-processing work

as well as `dials` and `tune`, which help manage and optimize tuning parameters in certain types of models.



# Today's Data (from Class 08)

## Heart and Estrogen/Progestin Study (HERS)

- Clinical trial of hormone therapy for the prevention of recurrent heart attacks and deaths among 2763 post-menopausal women with existing coronary heart disease (see Hulley et al 1998 and many subsequent references, including Vittinghoff, Chapter 4.)
- We're excluding the women in the trial with a diabetes diagnosis and those with missing LDL values.

```
hers_raw <- read_csv(here("data/hersdata.csv")) %>%  
  clean_names()
```

```
hers_new <- hers_raw %>%  
  filter(diabetes == "no") %>%  
  filter(complete.cases(ldl1, ldl)) %>%  
  select(subject, ldl1, ldl, age, ht, globrat)
```

## hers\_new Codebook (n = 1925)

Variable	Description
subject	subject code
ht	factor: hormone therapy or placebo
ldl	baseline LDL cholesterol in mg/dl
age	baseline age in years
globrat	baseline self-reported health (5 levels)
ldl1	LDL at first annual study visit
diabetes	yes or no (all are no in our sample)

**Goal** Predict percentage change in ldl from baseline to followup, using baseline age, ht, ldl and globrat, restricted to women without diabetes.


# Steps we'll describe today

- ➊ Create our outcome and consider a transformation.
- ➋ Split the data into training and testing samples.
- ➌ Build a recipe for our model.
  - Specify roles for outcome and predictors.
  - Deal with missing data in a reasonable way.
  - Complete all necessary pre-processing so we can fit models.
- ➍ Specify a modeling engine for each fit we will create.
  - There are five available engines just for linear regression!
- ➎ Create a workflow for each engine and fit model to the training data.
- ➏ Compare coefficients graphically from two modeling approaches.
- ➐ Assess performance in the models we create in the training data.
- ➑ Compare multiple models based on their performance in test data.

Key Reference: Kuhn and Silge, *Tidy Modeling with R* or TMWR

# Stage 1: Create our outcome

```
hers_new <- hers_new %>%  
  mutate(ldl_pch = 100*(ldl1 - ldl)/ldl)
```



min	Q1	median	Q3	max	mean	sd	n	missing
-80.9	-21	-8.9	5.6	137.4	-6.5	22.8	1925	0

## Stage 2: Creating Training and Test Samples



rsample

rsample provides infrastructure for efficient data splitting and resampling.

[Go to package ...](#)

Here, we'll use the rsample package to split our data.

```
set.seed(20210309)
hers_split <- initial_split(hers_new, prop = 0.8)

hers_train <- training(hers_split)
hers_test  <- testing(hers_split)
```

We start with 1925 women in `hers_new`, which we split into 1540 women in the training sample, leaving 385 women in the testing sample.

# What else can we do with `rsample`?

- Stratified sampling (splitting) on a categorical variable to ensure similar distributions of those categories in the training and testing groups.

```
initial_split(hers_new, prop = 0.8, strata = ht)
```

- What if you have time series data?
  - Use `initial_time_split()` to identify the first part of the data as the training set and the rest in the testing set; this assumes the data were pre-sorted in a sensible order.

The test set should **always** resemble new data that will be given to the model.

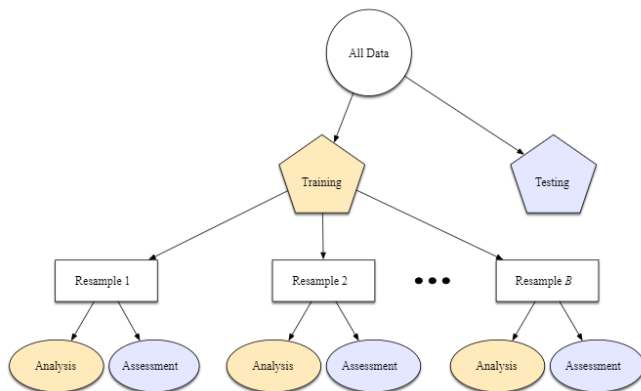
*A test set should be avoided only when the data are pathologically small.*

- TMWR, Section 5.2

# What about a validation set?

- Would like to avoid overfitting (where the models do much better on the training set samples than you do on the test set)
- Idea is to hold back a validation set of data to measure performance while training prior to moving on with a model to the test set.
- This is really just a special case of a resampling method used on the training set, as described in TMWR section 10 (see next slide).

# From TMWR, Section 10.2



Resampling is only conducted on the training set. The test set is not involved. For each iteration of resampling, the data are partitioned into two subsamples:

- The model is fit with the **analysis set**.
- The model is evaluated with the **assessment set**.



## Stage 3: Pre-Processing the Data



### recipes

recipes is a tidy interface to data pre-processing tools for feature engineering. [Go to package ...](#)

We'll build a **recipe** for our pre-modeling work. This might include:

- establishing the roles (outcome, predictors, identifiers) for variables
- pre-processing steps for predictors (feature engineering)
  - transforming predictors, including all of our usual power transformations, but also centering, scaling or normalizing and more complex mutations
  - creating dummy (indicator) variables for categorical data
  - dealing with factors and factor levels
  - including interactions, polynomials or splines
  - filtering out variables with zero variance
  - dealing with missing data via imputation or removal

<https://www.tidymodels.org/find/recipes/> lists all available recipes

# Building a Recipe for our modeling

```
hers_rec <-  
  recipe(ldl_pch ~ age + ht + ldl + globrat,  
          data = hers_new) %>%                # 1  
  step_bagimpute(all_predictors()) %>%         # 2  
  step_poly(ldl, degree = 2) %>%              # 3  
  step_dummy(all_nominal()) %>%               # 4  
  step_normalize(all_predictors())            # 5
```

Warning: `step\_bagimpute()` was deprecated in recipes 0.1.16.  
Please use `step\_impute\_bag()` instead.

This warning is displayed once every 8 hours.

Call `lifecycle::last\_lifecycle\_warnings()` to see where this

- 1 Specify the roles for the outcome and the predictors.
- 2 Impute missing predictors with bagged tree models.
- 3 Use an orthogonal polynomial of degree 2 with the baseline LDL data.
- 4 Form dummy variables to represent all categorical variables.
- 5 Normalize (subtract mean and divide by SD) all quantitative predictors.

# Column Roles

```
hers_rec <-  
  recipe(ldl_pch ~ age + ht + ldl + globrat,  
    data = hers_new)
```

- Everything to the left of the ~ is an outcome.
- Everything to the right of the ~ is a predictor.

Sometimes we want to assign other roles, like “id” for an important identifier that isn’t either a predictor or an outcome, or “split” for a splitting variable.

- Any character string can be a role, and columns can have multiple roles
- `add_role()`, `remove_role()` and `update_role()` functions are helpful

# Common steps used in building a recipe (1/5)

- Power Transformations of Predictors

- `step_log(x1, base = 10)` (default base is  $\exp(1)$ ), `step_sqrt`, `step_inverse`
- `step_BoxCox()` will transform predictors using a simple Box-Cox transformation to make them more symmetric (remember this does require a strictly positive variable, and will be something we'd use more for an outcome using the residuals for a statistical model).
- `step_YeoJohnson()` uses the Yeo\_Johnson transformation (again, typically on the outcome model) which is like Box-Cox but doesn't require the input variables to be strictly positive.

- `step_logit` and `step_invlogit`

- Non-Linear Terms for Quantitative Predictors

- `step_poly()` produces orthogonal polynomial basis functions
- `step_ns(x5, deg_free = 10)` from the `splines` package can create things called natural splines - the number of spline terms is a tuning parameter, `step_bs()` adds B-spline basis functions

# Common steps used in building a recipe (2/5)

- Dealing with Categorical Predictors

- `step_dummy(all_nominal())` which converts all factor or categorical variables into indicator (also called dummy) variables: numeric variables which take 1 and 0 as values to encode the categorical information
  - Other helpful selectors: `all_numeric()`, `all_predictors()` and `all_outcomes()`
  - If you want to select specific variables, you could use `step_dummy(x2, x3)`
- `step_relevel()` reorders the provided factor columns so that a level you specify is first (the baseline)
- If you have ordered factors in R, try `step_unorder()` to convert to regular factors or `step_ordinalscore()` to map specific numeric values to each factor level

# Common steps used in building a recipe (3/5)

- Dealing with Categorical Predictors (continued)
  - `step_unknown()` to change missing values in a categorical variable to a dedicated factor level
  - `step_novel()` creates a new factor level that may be encountered in future data
  - `step_other()` converts infrequent values to a catch-all labeled “Other” using a threshold
    - `step_other(x5, threshold = 0.05)` places bottom 5% of data in `x5` into “other”.
- Create Interaction Terms
  - `step_interact(~ interaction terms)` can be used to set up interactions
- Filter rows?
  - `step_filter()` can be used to filter rows using `dplyr` tools

# Common steps used in building a recipe (4/5)

- `step_mutate()` can be used to conduct a variety of basic operations
- `step_ratio()` can be used to create ratios of current variables
- Centering and Scaling Predictors
  - `step_normalize()` to center and scale quantitative predictors
  - `step_center()` just centers predictors
  - `step_scale()` just scales numeric data and
  - `step_range()` to scale numeric data to a specific range
- Zero Variance Filters
  - `step_zv()` is the zero variance filter which removes variables that contain only a single value.
  - `step_nzv()` removes variables with very few unique values or for whom the ratio of the frequency of the most common value to the second most common value is large

# Common steps used in building a recipe (5/5)

- Step options for imputation include things like
  - `step_meanimpute()` and `step_medianimpute()` to impute with mean or median,
  - `step_modelimpute()` to impute nominal data using the most common value,
  - `step_bagimpute()` for imputation via bagged trees,
  - `step_knnimpute()` to impute via k-nearest neighbors
- `step_naomit()` can be used to remove observations with missing values

<https://www.tidymodels.org/find/recipes/> lists all available recipes



## Stage 4: Specify `lm` modeling engine for `fit1`



### `parsnip`

`parsnip` is a tidy, unified interface to models that can be used to try a range of models without getting bogged down in the syntactical minutiae of the underlying packages. [Go to package ...](#)

```
hers_lm_model <- linear_reg() %>% set_engine("lm")
```

Other available engines for linear regression include:

- `stan` to fit Bayesian models
- `spark`
- `keras`

All `parsnip` models can be found at

<https://www.tidymodels.org/find/parsnip/>

## Stage 4: Specify stan modeling engine for fit2

As an alternative, we'll often consider a Bayesian linear regression model as fit with the “stan” engine. This requires the pre-specification of a prior distribution for the coefficients, for instance:

```
prior_dist_int <- rstanarm::student_t(df = 1)
prior_dist_preds <- rstanarm::normal(0, 5)

hers_stan_model <- linear_reg() %>%
  set_engine("stan",
             prior_intercept = prior_dist_int,
             prior = prior_dist_preds)
```

## Stage 5: Create a workflow for the `lm` model



### workflows

workflows bundle your pre-processing, modeling, and post-processing together. [Go to package ...](#)

```
hers_lm_wf <- workflow() %>%  
  add_model(hers_lm_model) %>%  
  add_recipe(hers_rec)
```

### Fit the `lm` model to the training sample

```
fit1 <- fit(hers_lm_wf, hers_train)
```

We'll show the `fit1` results on the next slide.

```

> fit1
== Workflow [trained] =====
Preprocessor: Recipe
Model: linear_reg()

-- Preprocessor -----
4 Recipe Steps

* step_bagimpute()
* step_poly()
* step_dummy()
* step_normalize()

-- Model -----

Call:
stats::lm(formula = ..y ~ ., data = data)

Coefficients:
      (Intercept)          age      1d1_poly_1      1d1_poly_2
      -6.0248        -1.6396        -8.0728         2.5596
      ht_placebo      globrat_fair      globrat_good      globrat_poor
       5.3921        -1.3379        -1.8050        -0.7685
globrat_very.good
      -1.4063

```

# Tidy the coefficients for fit1?



## broom

broom converts the information in common statistical R objects into user-friendly, predictable formats. [Go to package ...](#)

term	estimate	std.error	conf.low	conf.high
(Intercept)	-6.368	0.523	-7.394	-5.342
age	-0.772	0.525	-1.802	0.258
ldl_poly_1	-7.857	0.524	-8.884	-6.829
ldl_poly_2	2.236	0.524	1.208	3.265
ht_placebo	4.893	0.523	3.866	5.920
globrat_fair	-0.723	1.002	-2.689	1.242
globrat_good	-1.569	1.196	-3.915	0.777
globrat_poor	-1.123	0.572	-2.244	-0.001
globrat_very.good	-1.390	1.114	-3.574	0.795

# Want to glance at the fit1 summaries?

```
fit1 %>% extract_fit_parsnip() %>%  
  glance() %>% select(1:6) %>% kable(dig = 3)
```

r.squared	adj.r.squared	sigma	statistic	p.value	df
0.18	0.175	20.522	41.923	0	8

```
fit1 %>% extract_fit_parsnip() %>%  
  glance() %>% select(7:12) %>% kable(dig = 1)
```

logLik	AIC	BIC	deviance	df.residual	nobs
-6833.8	13687.6	13741	644801.3	1531	1540

## Stage 5: Create a workflow for the stan model

```
hers_stan_wf <- workflow() %>%  
  add_model(hers_stan_model) %>%  
  add_recipe(hers_rec)
```

### Fit the stan model to the training sample

```
set.seed(43202)  
fit2 <- fit(hers_stan_wf, hers_train)
```

We'll show the fit2 results on the next slide.

```

> fit2
== Workflow [trained] =====
Preprocessor: Recipe
Model: linear_reg()

-- Preprocessor -----
4 Recipe Steps

* step_bagimpute()
* step_poly()
* step_dummy()
* step_normalize()

-- Model -----
stan_glm
family:      gaussian [identity]
formula:     ..y ~ .
observations: 1444
predictors:  9
-----
              Median MAD_SD
(Intercept)   -5.9    0.6
age            -1.6    0.5
ldl_poly_1     -8.0    0.6
ldl_poly_2      2.5    0.6
ht_placebo      5.3    0.6
globrat_fair    -1.1    1.0
globrat_good   -1.5    1.1
globrat_poor    -0.7    0.6
globrat_very.good -1.1    1.1

Auxiliary parameter(s):
      Median MAD_SD
sigma 20.8    0.4
-----
* For help interpreting the printed output see ?print.stanreg
* For info on the priors used see ?prior_summary.stanreg

```



# Tidy the fit2 coefficients?

The stan model requires the broom.mixed package to tidy the fit.

```
broom.mixed::tidy(fit2, conf.int = T) %>% kable(dig = 3)
```

term	estimate	std.error	conf.low	conf.high
(Intercept)	-6.286	0.521	-7.150	-5.445
age	-0.768	0.515	-1.579	0.104
ldl_poly_1	-7.761	0.527	-8.627	-6.902
ldl_poly_2	2.219	0.514	1.374	3.112
ht_placebo	4.833	0.540	3.989	5.694
globrat_fair	-0.634	0.933	-2.161	0.931
globrat_good	-1.407	1.112	-3.260	0.397
globrat_poor	-1.081	0.571	-1.988	-0.179
globrat_very.good	-1.258	1.077	-2.969	0.482

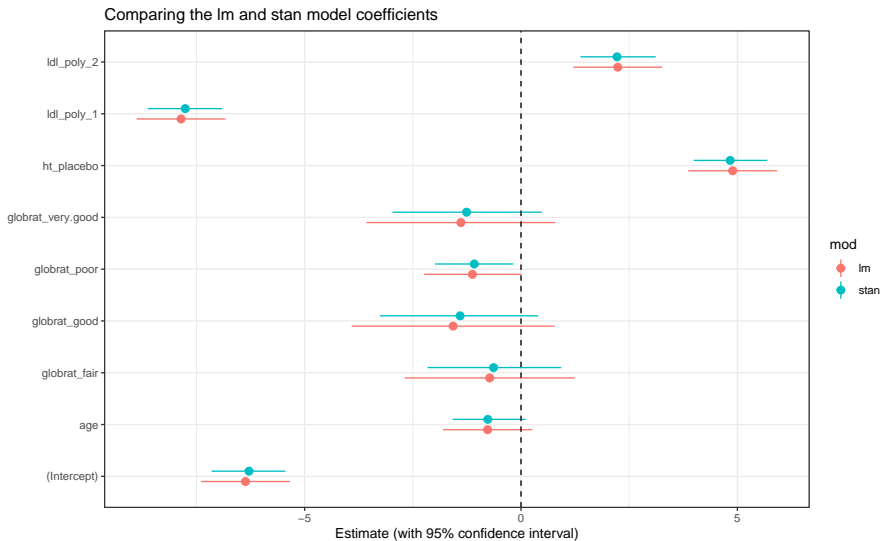
## Stage 6: Compare the coefficients of the fits

```
coefs_lm <- tidy(fit1, conf.int = TRUE) %>%  
  select(term, estimate, conf.low, conf.high) %>%  
  mutate(mod = "lm")  
  
coefs_stan <- tidy(fit2, conf.int = TRUE) %>%  
  select(term, estimate, conf.low, conf.high) %>%  
  mutate(mod = "stan")  
  
coefs_comp <- bind_rows(coefs_lm, coefs_stan)
```

# Graph the coefficients from the two models

```
ggplot(coefs_comp, aes(x = term, y = estimate, col = mod,  
                      ymin = conf.low, ymax = conf.high)) +  
  geom_point(position = position_dodge2(width = 0.4)) +  
  geom_pointrange(position = position_dodge2(width = 0.4)) +  
  geom_hline(yintercept = 0, lty = "dashed") +  
  coord_flip() +  
  labs(x = "", y = "Estimate (with 95% confidence interval)",  
       title = "Comparing the lm and stan model coefficients")
```

# Graph the coefficients from the two models



# Stage 7. Assess performance in the training data



yardstick

yardstick measures the effectiveness of models using performance metrics.

[Go to package ...](#)

Available regression performance metrics include:

- `rsq` (r-squared, via correlation - always between 0 and 1)
- `rmse` (root mean squared error)
- `mae` (mean absolute error)
- `rsq_trad` (r-squared, calculated via sum of squares)

but there are many, many more. Let's select two...

```
mets <- metric_set(rsq, rmse)
```

# Make predictions using fit1 in training sample

```
lm_pred_train <-  
  predict(fit1, hers_train) %>%  
  bind_cols(hers_train %>% dplyr::select(ldl_pch))  
  
# remember  
mets <- metric_set(rsq, rmse)  
  
lm_res_train <-  
  mets(lm_pred_train, truth = ldl_pch, estimate = .pred)
```

We'll see the results in a moment.

# Make predictions using fit2 in training sample

```
stan_pred_train <-  
  predict(fit2, hers_train) %>%  
  bind_cols(hers_train %>% select(ldl_pch))  
  
# remember  
mets <- metric_set(rsq, rmse)  
  
stan_res_train <-  
  mets(stan_pred_train, truth = ldl_pch, estimate = .pred)
```

We'll see the results from each fit on the next slide.

# fit1 and fit2 performance in the training sample

from fit1 with lm:

```
lm_res_train %>% kable()
```

.metric	.estimator	.estimate
rsq	standard	0.1796985
rmse	standard	20.4622118

from fit2 with stan:

```
stan_res_train %>% kable()
```

.metric	.estimator	.estimate
rsq	standard	0.1796909
rmse	standard	20.4626978



# What about adjusted $R^2$ ?

The `yardstick` package doesn't use adjusted  $R^2$ .

- `tidymodels` wants you to compute performance on a separate data set for comparing models rather than doing what adjusted  $R^2$  tries to do, which is evaluate the model on the same data as were used to fit the model.

## Stage 8. Compare model performance on test data

```
lm_pred_test <-  
  predict(fit1, hers_test) %>%  
  bind_cols(hers_test %>% dplyr::select(ldl_pch))  
  
lm_res_test <-  
  mets(lm_pred_test, truth = ldl_pch, estimate = .pred)  
  
stan_pred_test <-  
  predict(fit2, hers_test) %>%  
  bind_cols(hers_test %>% select(ldl_pch))  
  
stan_res_test <-  
  mets(stan_pred_test, truth = ldl_pch, estimate = .pred)
```

# fit1 and fit2 performance in the test sample

from fit1 with lm:

```
lm_res_test %>% kable()
```

.metric	.estimator	.estimate
rsq	standard	0.199772
rmse	standard	21.082747

from fit2 with stan:

```
stan_res_test %>% kable()
```

.metric	.estimator	.estimate
rsq	standard	0.1997987
rmse	standard	21.0858904

# Where to Learn More

- Tidy Modeling with R by Max Kuhn and Julia Silge.
  - The Basics section (Chapters 4-9) as well as chapters 10-11 were my main tools for learning about these ideas.
- Julia Silge has many nice videos on YouTube demonstrating various things that `tidymodels` can accomplish.
  - I've recommended several in the Class 10 README.
- Lab 3 Part B requires you to use `tidymodels` approaches to complete a linear regression model using two different fitting engines

## Next Time

We'll apply ideas from the `tidymodels` framework to fit a logistic regression model.